

RAISR

Rapid and Accurate Image Super Resolution



Nearest Neighbor



Bilinear



Nearest Neighbor



Bilinear



Cubic



Lanczos4



Cubic



Lanczos4

hqx -- Pixel Art Scaling

- Intended for simple 2D art



Nearest Neighbor

- Use patch around each pixel
- Neighbors are “close” or “distant”
- Lookup neighborhood in LUT
- Determine filter



hq3x

Not intended for photographs!



Not intended for photographs!



Nearest Neighbor



hq4x

Example-Based Approaches

- Build a dictionary from low resolution (LR) patches to high resolution (HR) patches
 - Dictionary is very large, so sparse coding is typically used
 - Usually limited to fixed upscaling factors
- Anchored Neighborhood Regression (ANR) and A+
 - Uses neighbor embedding and ridge regression
 - Somewhat fast
- SRCNN
 - End-to-end approach with a convolutional neural network (CNN)
 - Slow



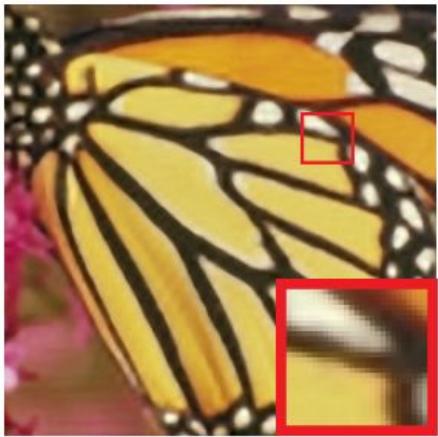
(a) High-res



(b) A+



(c) SRCNN



(d) JSB-NE



(e) CSCN



(f) DEGREE

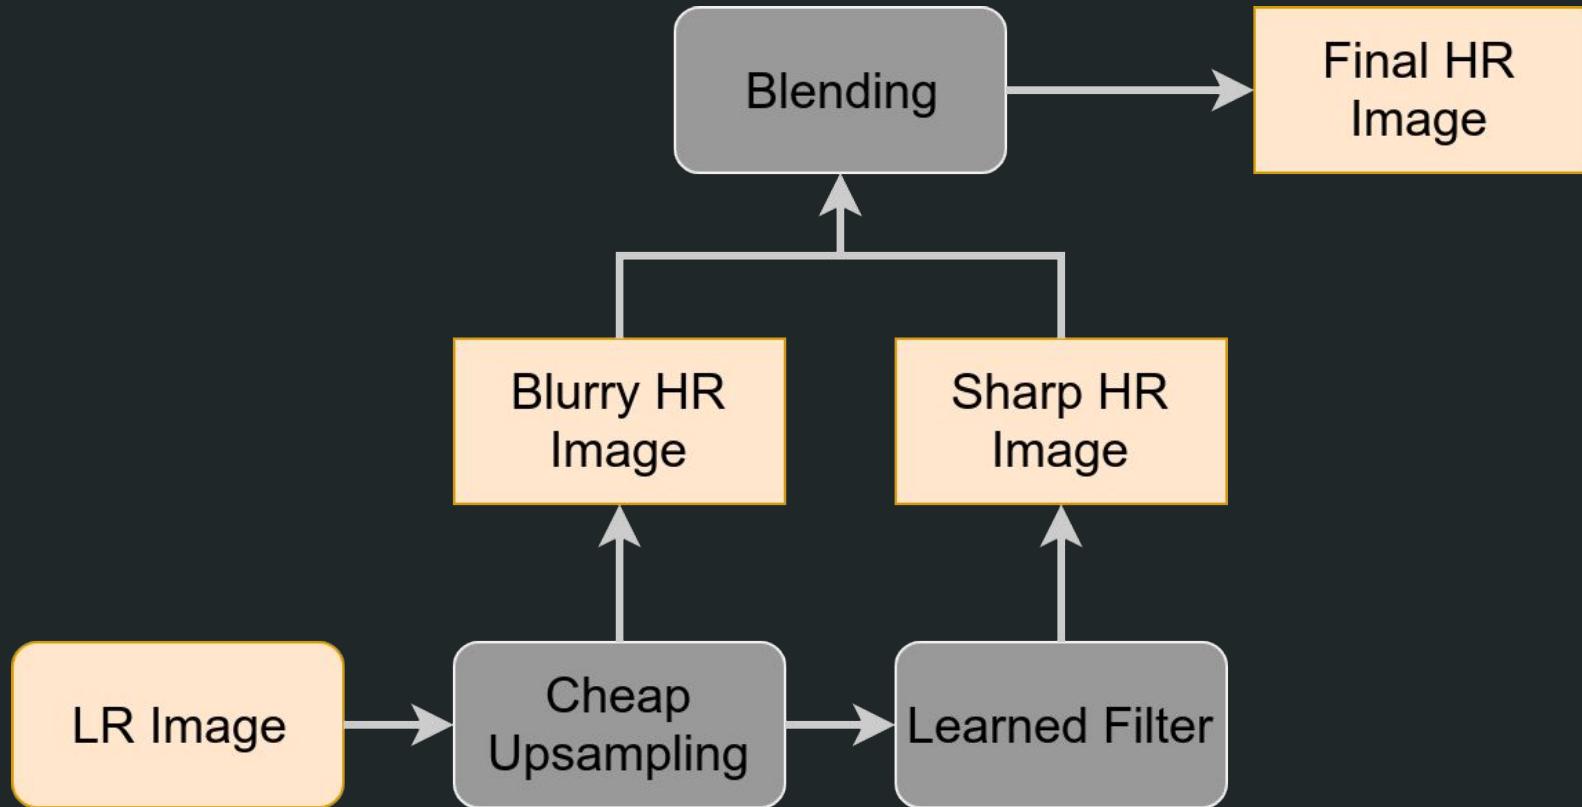
RAISR -- History

- Developed by Google; paper published in 2017
- Intended for mobile devices
- Lightweight, easy to parallelize, low memory requirements
- Performs roughly as well as SRCNN

RAISR -- Overview

- Use a cheap filter to scale the image up
- Sharpen the image using pre-learned filters
 - The filter to use depends on the neighborhood of the pixel
- Perform a blending pass to remove undesirable oversharpening effects

RAISR -- Overview

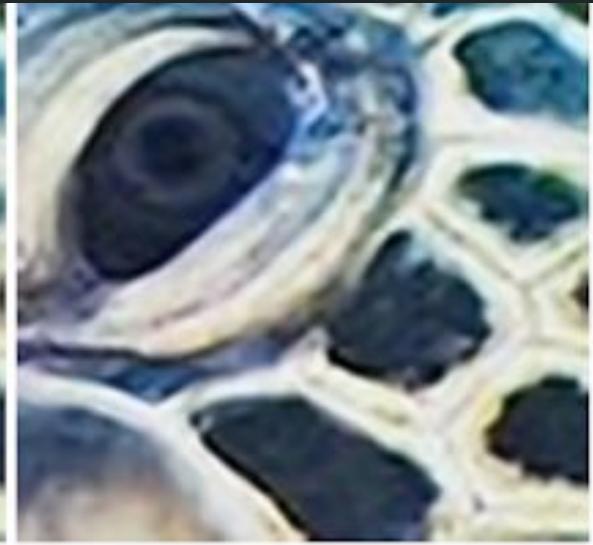




Nearest Neighbor

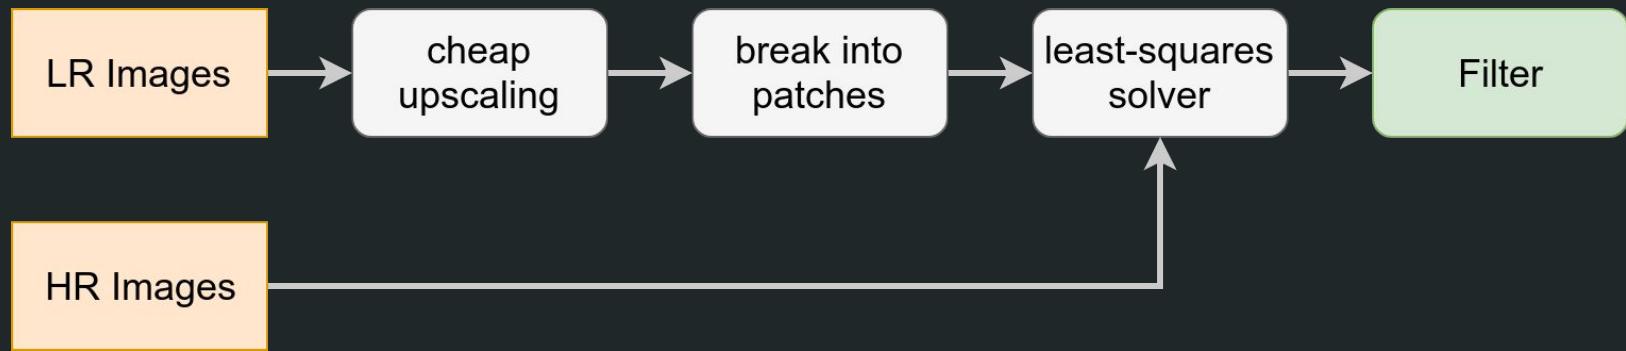


Bicubic Upsampling



RAISR Upsampling

Learning Phase



Objective Function

- Goal: A database of filters that minimizes the Euclidean distance between the **generated HR images** x_i and the **training HR images** y_i .
- Minimize

$$\sum_{i=1}^L \left\| \begin{bmatrix} A_i | h \end{bmatrix} - \begin{bmatrix} b_i \end{bmatrix} \right\|_2^2$$

All d by d patches in y_i
[$M N \times d^2$]

d by d filter to learn
[$d \times d$]

All pixels in x_i , correspond to center of y_i patches
[$M N$]

Actual Objective Function

- A is a big matrix and we do not need all of the patches
- Much simpler computational/memory requirements in this form
- Minimize

$$\|Qh - V\|_2^2$$

$Q = A^T A$
[$d^2 \times d^2$]

d by d filter to learn
[$d \times d$]

$V = A^T b$
[d^2]

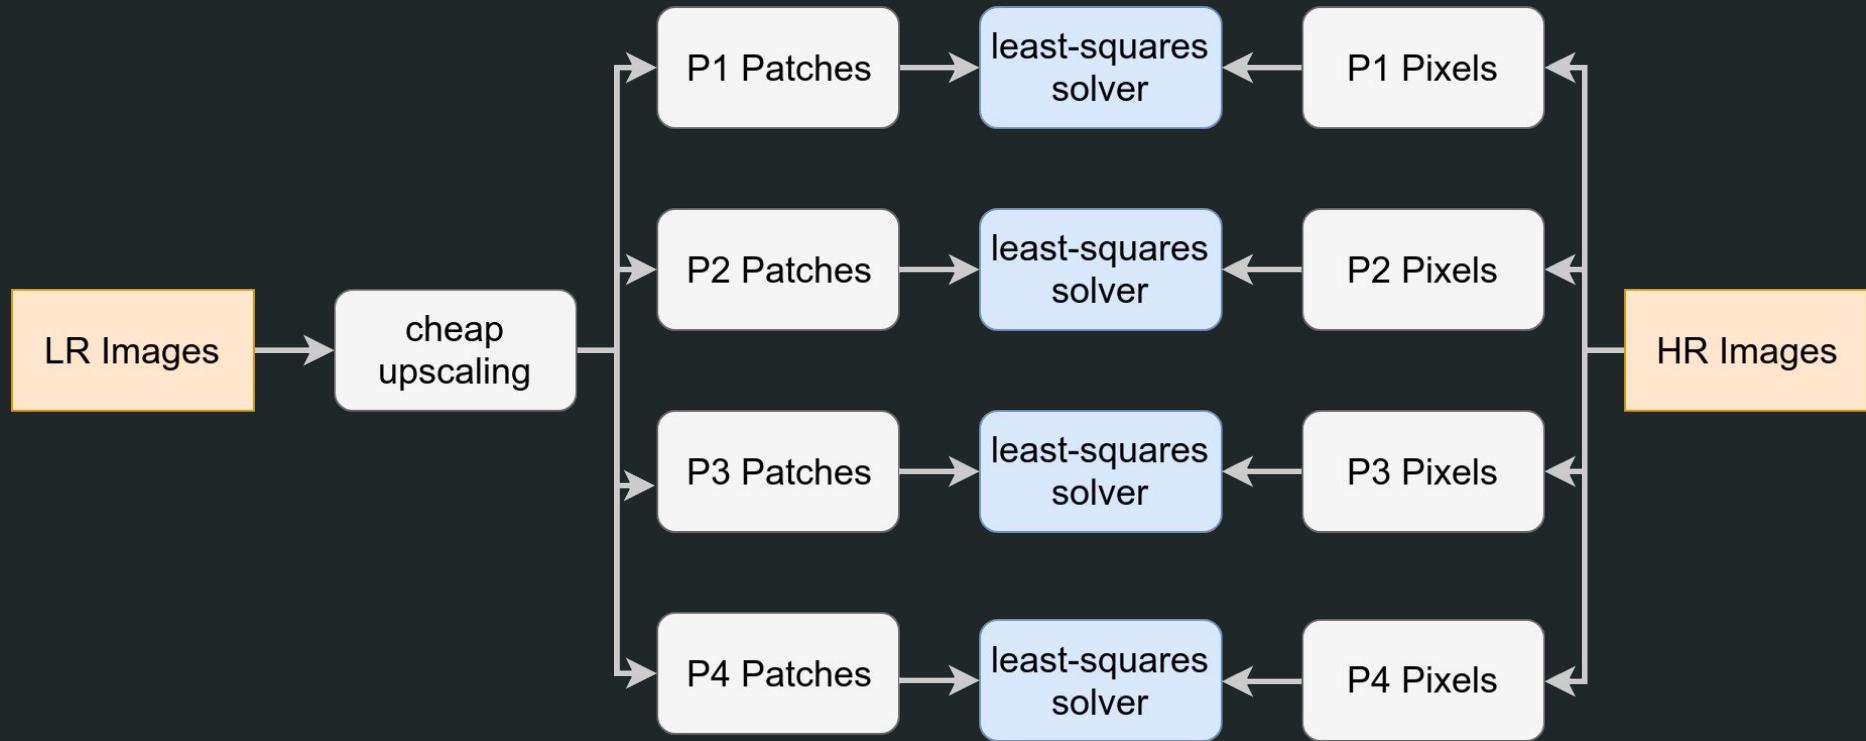
- Q is small!
- We don't need to keep the whole image in memory

Shift-invariance and Aliasing

- The cheap upscaler is not shift invariant
- 2x upsample mean 4 types of pixels; 3x upsample means 9 types of pixels; etc
- Separate filter for different types of pixel

P1	P2	P1	P2	P1
P3	P4	P3	P4	P3
P1	P2	P1	P2	P1
P3	P4	P3	P4	P3
P1	P2	P1	P2	P1

Learning Phase (with Shift Invariance)



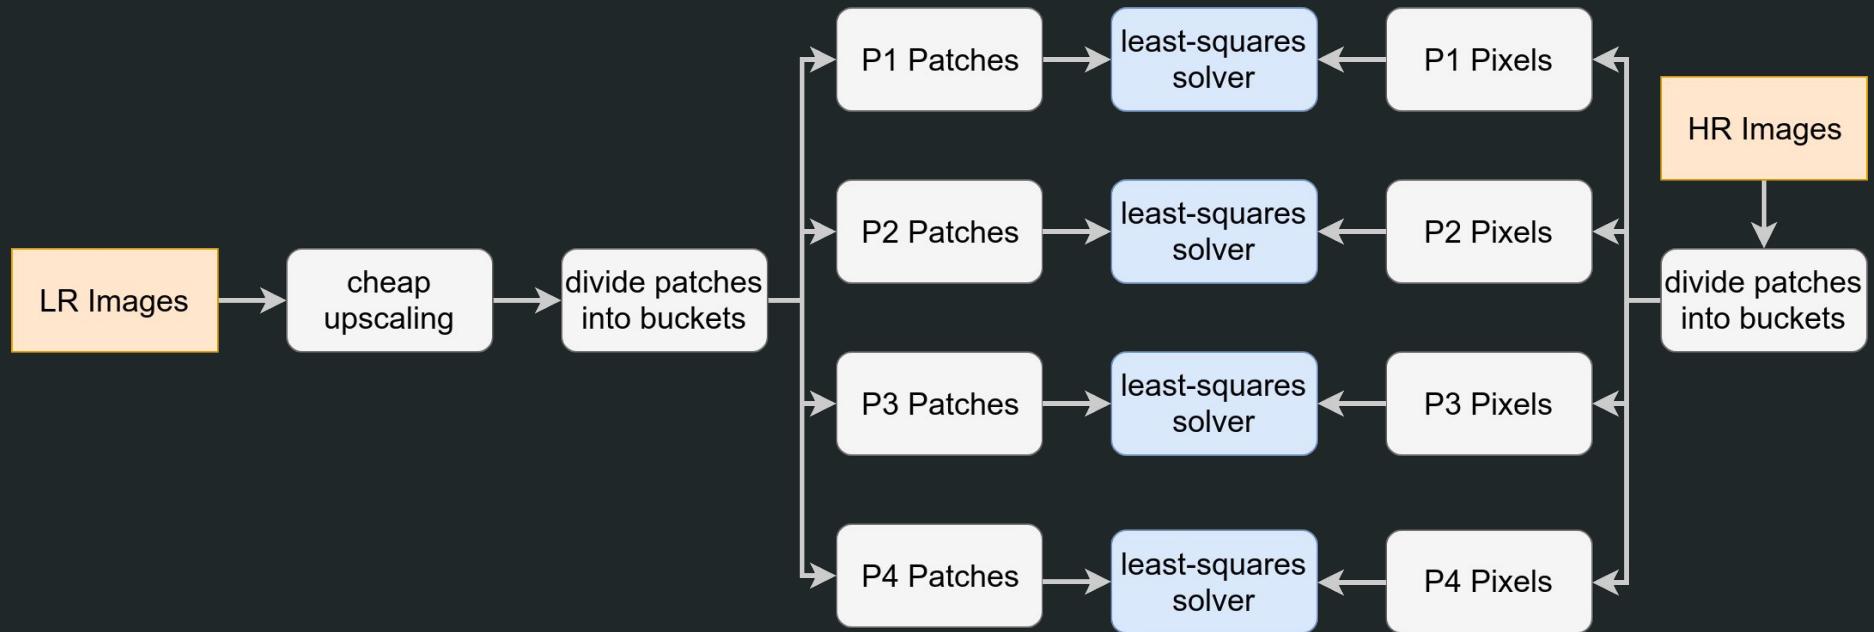
Hashing-Based Learning

- Global image filtering is fast, but not very accurate
- Operate on the **patch** level to improve performance
- Instead of clustering, we **hash** the patch to get a bucket
- Now we compute each pixel filter for each bucket

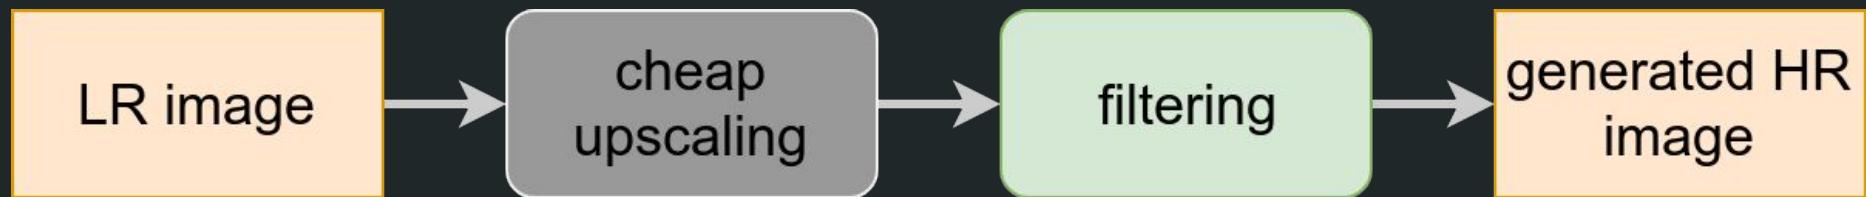
Hashing-Based Learning

- A 9x9 neighborhood of the surrounding pixels is the hash key
- $\text{Hash}(\text{Neighborhood}) = (\theta, \lambda, \mu)$
- Patches are divided into buckets; solver operates on individual buckets

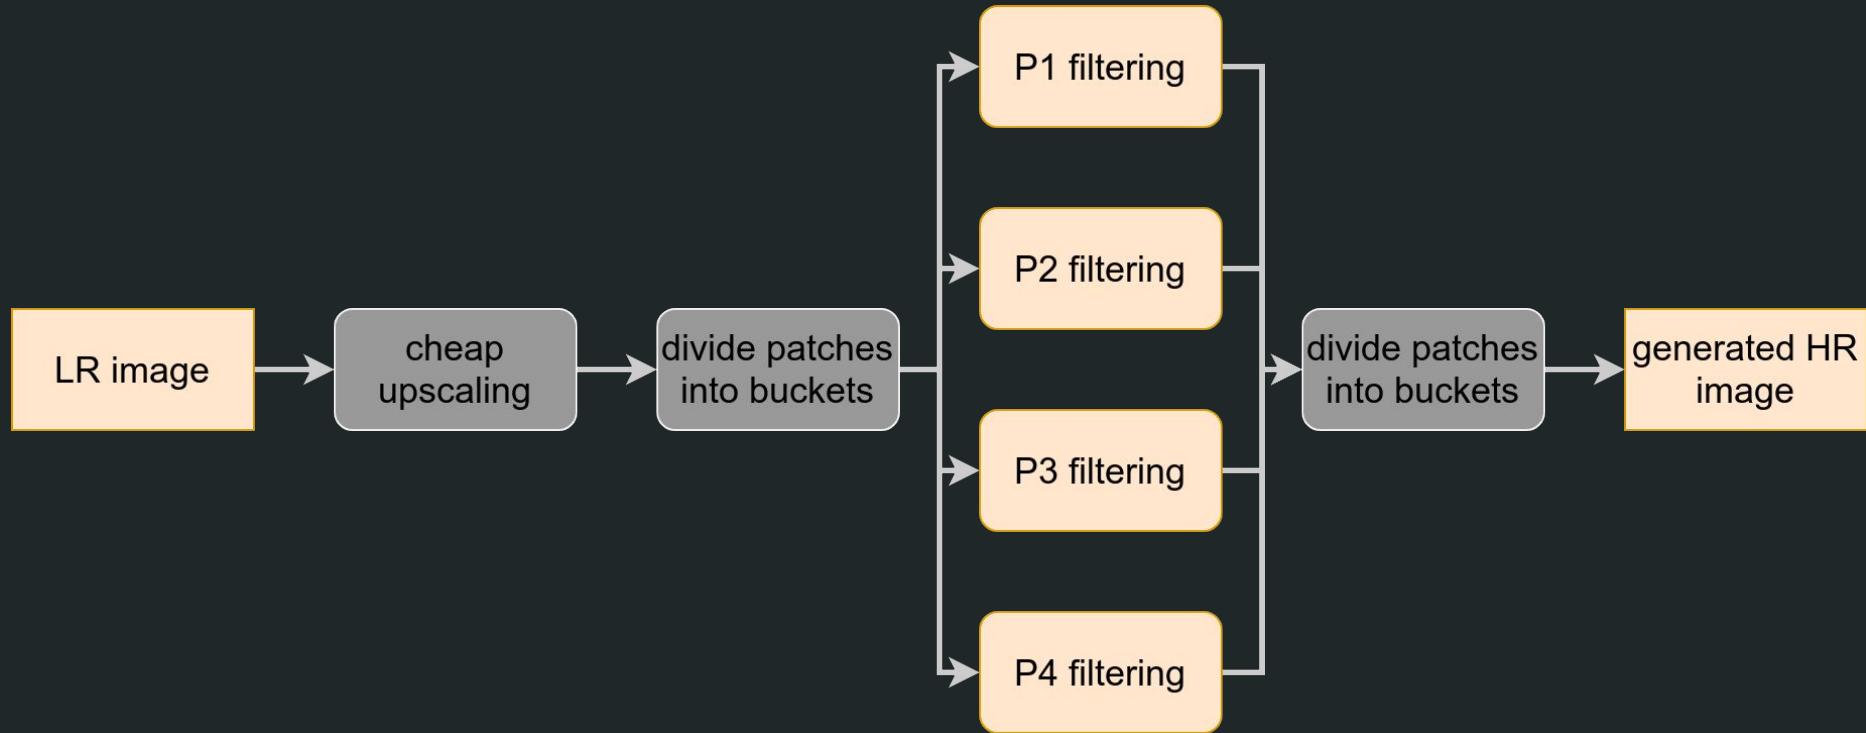
Learning Phase (Hash-based with Shift Invariance)



Inference Phase



Inference Phase (Hash-based with Shift Invariance)



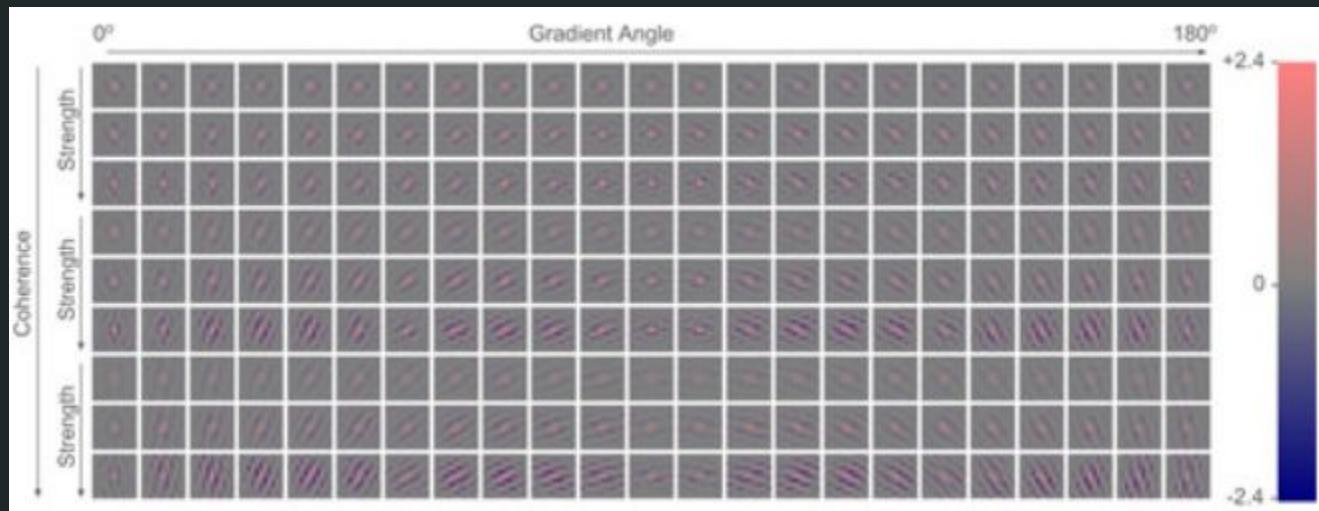
Inference Phase

- Hash the neighborhood
 - Find the bucket for that hash value
 - Select the correct pixel filter from the bucket
 - Apply the filter
-
- How does the hashing algorithm work?

Hashing algorithm

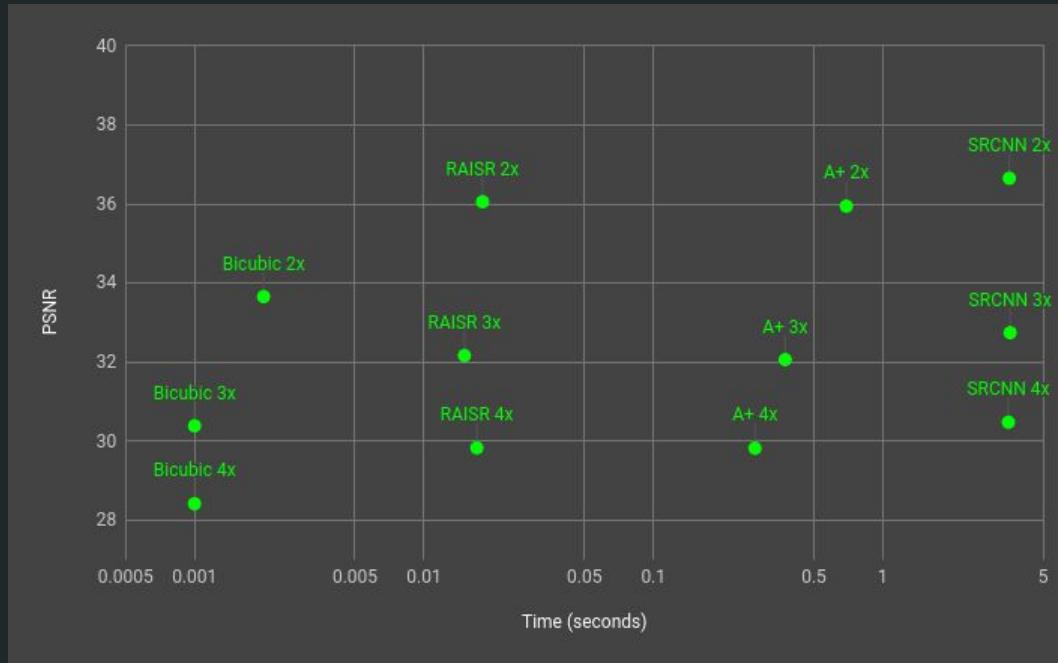
1. Collect a 9x9 neighborhood of the pixel
2. Compute the gradient of this patch
3. G is a $[n^2 \times 2]$ matrix where the rows are the gradients at each location
4. Compute $G^T W G$ where W is a normalized Gaussian kernel (applied separably)
 - a. Note that this is a $[2 \times 2]$ matrix
5. Compute the SVD of this matrix
6. The **angle (θ)** parameter is the angle of the eigenvector corresponding to the larger singular value
7. The **strength (λ)** parameter is the square root of the larger eigenvalue.
8. The **coherence (μ)** (strength + spread) is $\frac{\sqrt{\lambda_1} - \sqrt{\lambda_2}}{\sqrt{\lambda_1} + \sqrt{\lambda_2}}$
9. Quantize gradient statistics to buckets

Learned 2x Filters

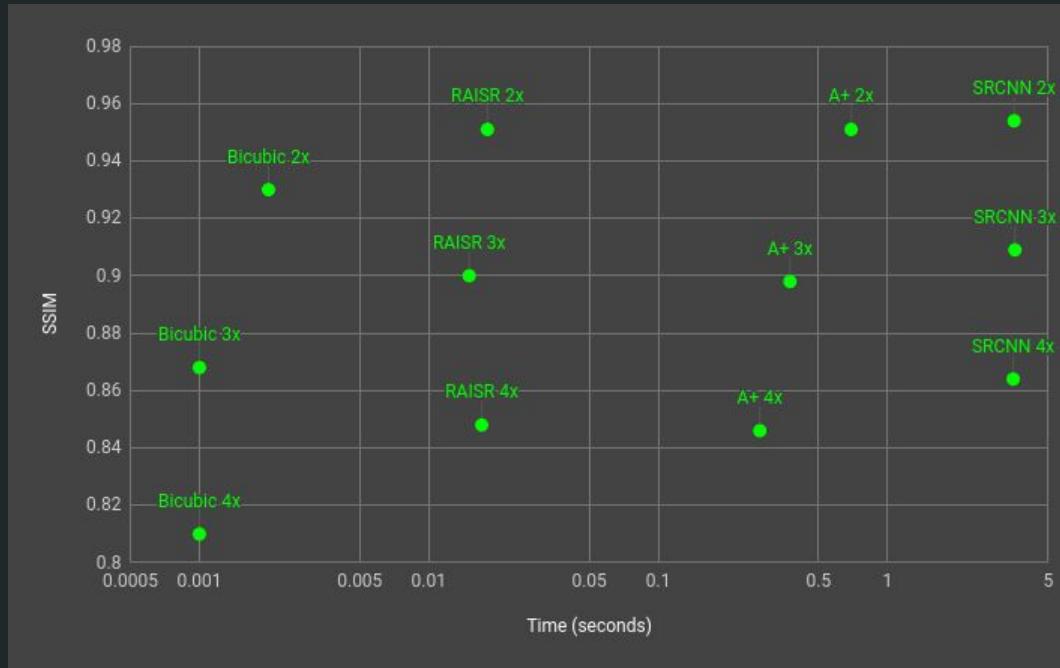


(Almost) Final Steps

1. $\text{FilterDatabase}[\theta, \lambda, \mu, \text{PixelType}] = h$
 - a. In the paper, filters were [11 x 11]
2. Apply the filter to the [11x11] neighborhood of the pixel
3. Add the (cheaply-scaled) chroma back in



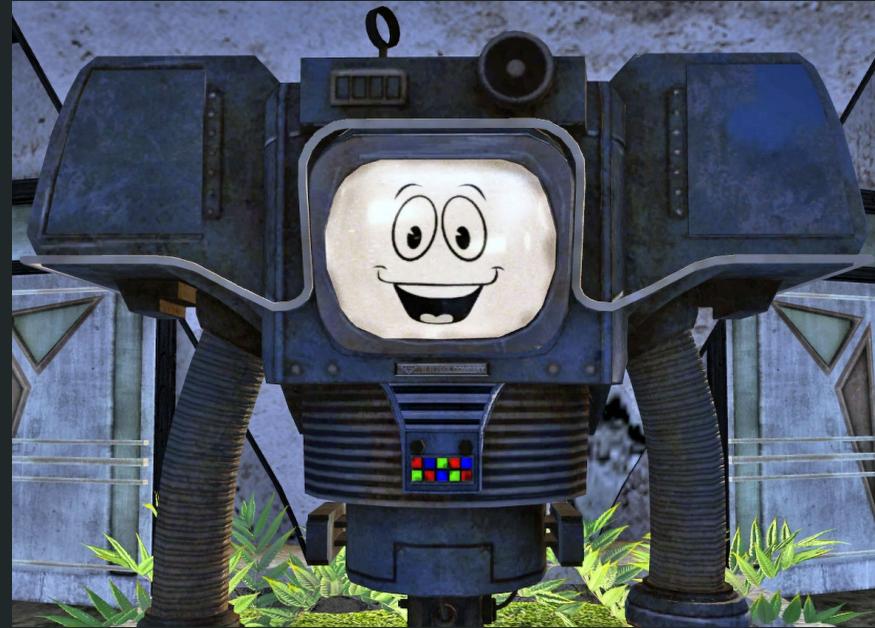
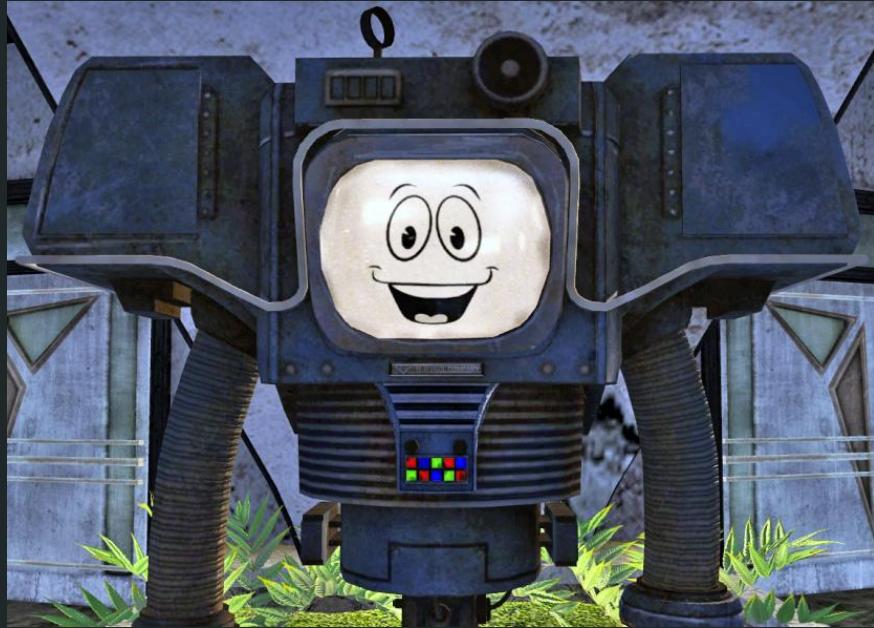
Evaluation (PSNR) -- 3.4GHz 6-Core Xeon



Evaluation (SSIM) -- 3.4GHz 6-Core Xeon

















Removing Artifacts

- So far RAISR worsens existing artifacts in image
- LR images are often distorted by an unknown PSF (e.g blur) or have JPEG compression artifacts
- Solution
 - To avoid oversharpening, train on LR images with **compression artifacts** to normal HR images
 - To avoid blurriness, train on normal LR images to **oversharpened** HR images.

Removing Artifacts



Bilinear



Uncompressed LR



Compressed LR

Avoiding Oversharpening

- Sadly this technique tends to amplify noise and create halos
- We want a cheap filter for **low**-frequency regions, learned filter for **high**-frequency regions
- Could cluster and handle areas differently, but that would be slow

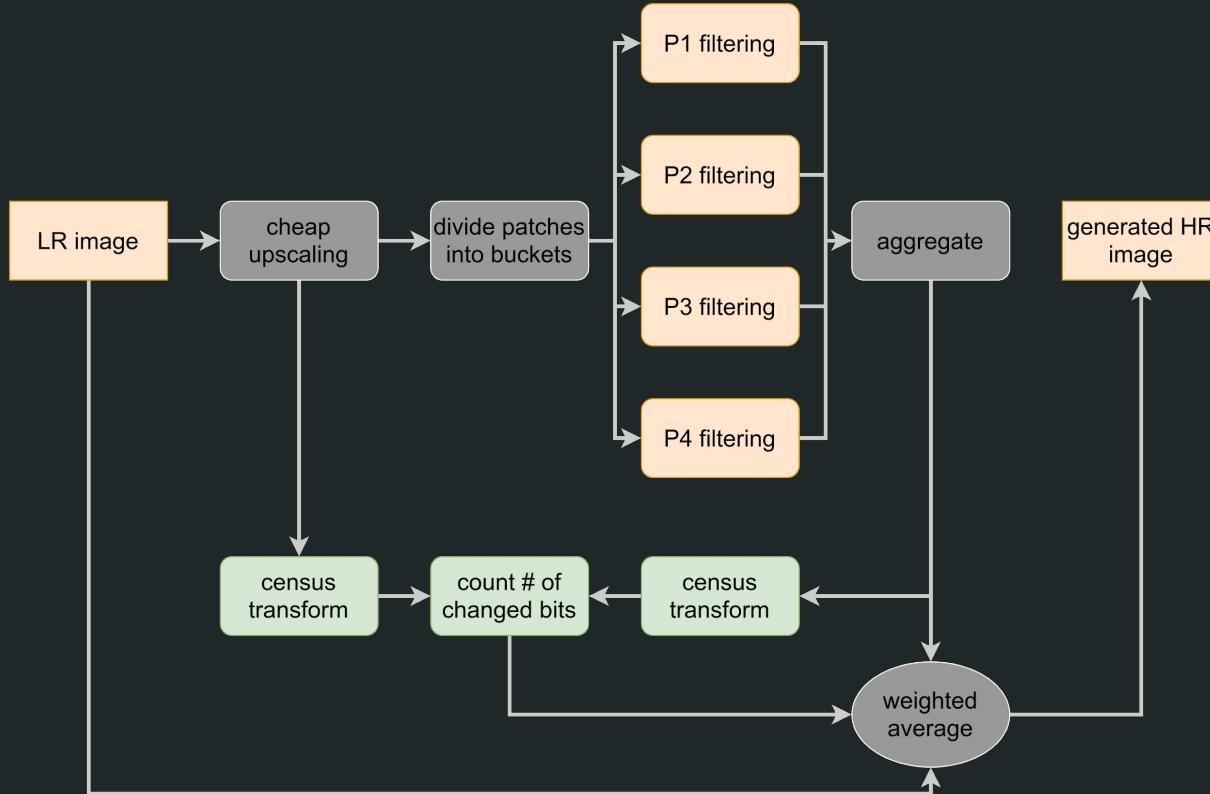
Contrast Enhancement

- Need a measure of structure
- Least Connected Component (LCC) of the Census Transform (CT) could be used to detect high-frequency regions
- We can either selectively amplify high frequencies, or enhance a variety of frequencies

Evaluating Local Change in Structure

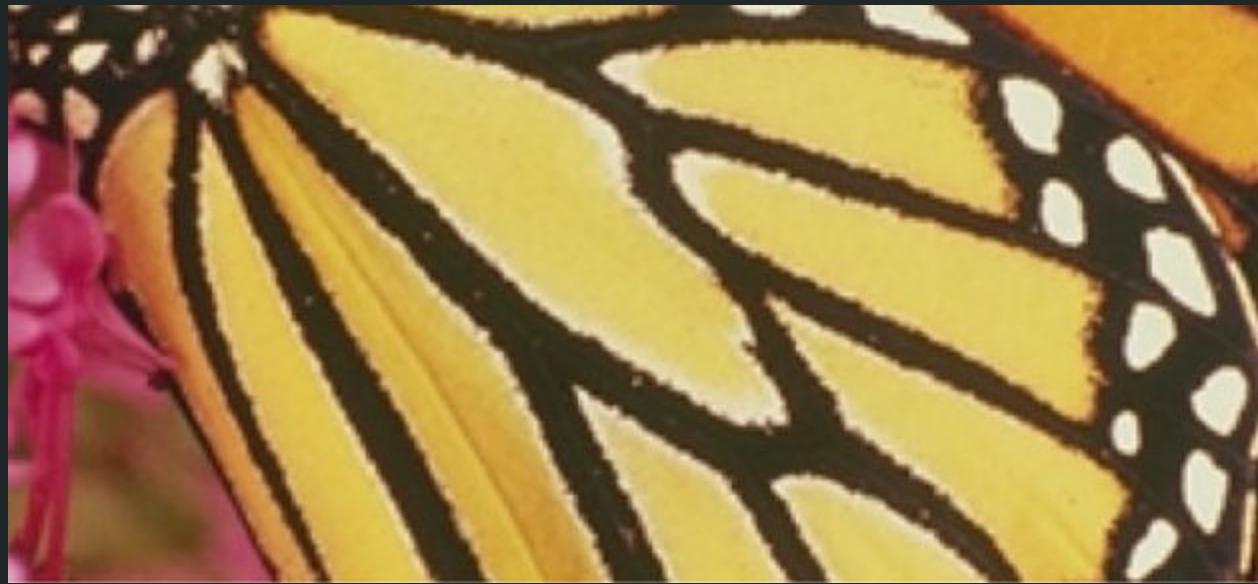
- Compute the CT of the cheap upscaled image and the filtered image
- Per pixel, evaluate the Hamming distance
- Distance is proportional to the change in structure
- This weight is used for blending both images together

Inference Phase (Everything + CE)

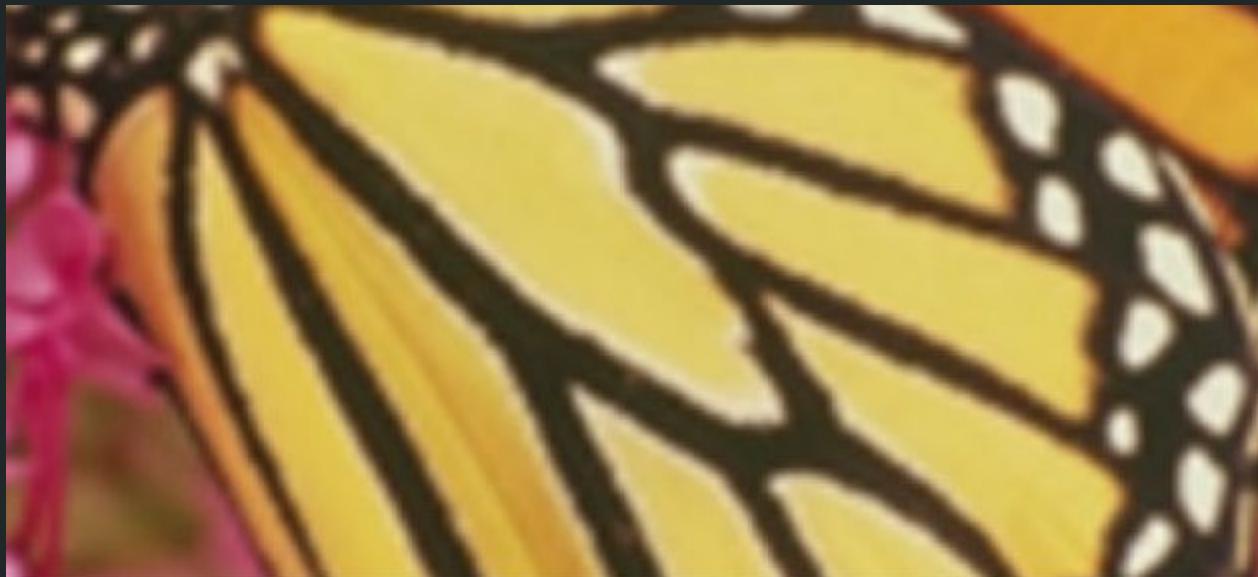


Pre-processing for Contrast Enhancement

- Use a **DoG** sharpener on the HR **training** images to increase **contrast**
- Automatically enhances details (HF) and contrast (MF/LF)



Original Image (HR)



Bicubic upscaling (LR \rightarrow HR)



Filtered (**not trained** on sharpened HR)



Filtered (**trained** on sharpened HR)



Blended (with bicubic)



Original Image (once again)

RAISR

- Lightweight
- Simpler and faster than Deep Learning techniques
- Works well for a variety of content -- games, photographs, video, etc

Some ideas

- Applying the [11x11] filter is likely the most expensive operation (~121 MADs)
 - Why not use separable convolutions (~22 MADs)?
- 3D applications may have access to the **depth buffer**
 - Could be used in the hashing algorithm or the filter